



A Multi-Method Content-Based Recommendation System for Netflix Data

DHEER GUPTA
VYAAS SUBRAMANIAN
MUHAMMAD YAHYA
MARCH, 2025

Abstract

Recommender systems have become integral to online platforms such as Netflix, Amazon Prime, and Hulu, providing tailored suggestions to users from extensive content catalogs. While collaborative filtering often headlines the research in recommender systems, content-based approaches remain vital in situations with limited user-user overlap or when item features can strongly signal user preferences. Recent developments suggest combining multiple content-based signals—such as item similarity, classification-based “like” prediction, clustering, and association rule mining—can yield a more robust personalized experience.

Hence, our research question is: *How can we develop a multi-faceted, purely content-based recommendation system for Netflix data that effectively leverages user likes/dislikes and advanced methods (kNN, Naïve Bayes, clustering, and association rules) to produce high-quality suggestions?*

1 Introduction

Netflix, like many other streaming services, provides basic recommendations based on algorithms that primarily focus on genre, viewing history, or user ratings. However, these systems tend to offer limited and often generalized suggestions that may not accurately reflect the unique tastes of each user. Users often find themselves scrolling endlessly through an overwhelming selection of content, unsure of which movies or TV shows to watch next. This issue is further compounded by the fact that most recommendation systems fail to account for subtle and diverse factors, such as specific genres, IMDB ratings, actors, or production companies, that could significantly enhance the personalization of the recommendations. The importance of an advanced recommendation system is clear, especially as competition within the streaming industry increases. A more personalized and effective recommendation engine could not only enhance user satisfaction but also improve engagement and retention for platforms like Netflix. Previous research in the field of recommendation systems has highlighted the importance of using more than just basic genre filtering to make personalized suggestions. For instance, hybrid approaches that combine collaborative filtering (which uses user behavior data) with content-based filtering (which uses item attributes like genre or director) have shown improved performance in generating relevant recommendations. Despite this, the lack of deep personalization based on granular user preferences remains a challenge. This paper seeks to address that challenge by developing a recommendation engine that leverages both user input and detailed content attributes such as genre, IMDB rating, and production company to

generate tailored suggestions. Specifically, the research question guiding this study. To answer our question, we have developed a more advanced recommendation engine using a dataset from Netflix, which includes a variety of content, such as movies, TV shows, and anime, from numerous genres. The dataset offers detailed content attributes, including genre, IMDB ratings, release year, and more, which provide the necessary information for making tailored recommendations. To gather additional user preferences, a Google form was created, allowing users to specify their likes and dislikes, genre preferences, and minimum IMDB rating. The study applies machine learning techniques, including cosine similarity, k-Nearest Neighbors (kNN), and ensemble methods, to predict the most relevant content based on both user input and content features. The findings from this research reveal that incorporating both granular user preferences and detailed content attributes allows the recommendation engine to generate more personalized and accurate suggestions compared to traditional models. Ultimately, the results show that a deeper level of personalization—considering specific tastes, ratings, and genres—leads to a significantly enhanced user experience. This paper concludes that future recommendation systems must incorporate these nuanced factors to create more engaging and user-centric platforms.

2 Literature Review

2.1 Early Recommendation Systems

Initial recommender systems leaned heavily on collaborative filtering (1; 2), deriving suggestions from user-user similarities. However, collaborative filtering struggles with the “cold start” problem (3) when new users or items have insufficient interactions. Meanwhile, content-based approaches (4; 5) overcame this by matching item attribute vectors to each user’s profile of previously liked items.

2.2 Hybrid & Multi-Signal Approaches

As recommendation research evolved, hybrid methods (6; 7) combined both collaborative and content-based, or integrated multiple content-based signals. Some authors introduced classification models (Naïve Bayes, SVM) to label items as “likely liked,” e.g. (8). Clustering such as K-Means or hierarchical clustering helps discover latent item groups (9). kNN item retrieval (10) can quickly find top local neighbors. Additionally, association rule mining (11) has proven valuable in discovering frequent co-attributes within item metadata.

2.3 Association Rule Mining for Content

Apriori, initially for market basket data (11), can detect frequent itemset patterns. In a recommendation context, each item’s attribute set (particularly genres) can be treated as a transaction (12). Identifying co-occurring attributes (“Independent Movies” \rightarrow “Dramas”) can refine the recommendation process by giving these patterns a small “bonus” during ranking.

2.4 Our Position

Motivated by these approaches, we propose a purely content-based system merging multiple signals:

- Cosine-based content similarity,
- kNN retrieval for local neighbors,
- Naïve Bayes classification for synthetic likes,
- K-Means clustering for user cluster preference,
- Apriori association rules for discovering genre synergy,

all integrated in a final ensemble. While we do not incorporate user-user collaborative data, we mitigate cold-start issues and generate robust, varied recommendations by fusing these distinct content methods.

3 Data and Analysis

3.1 Dataset Collection

We used a Netflix dataset from Kaggle (13) of approximately 8,797 entries. Each row is labeled “Movie” or “TV Show” and includes:

- Basic Info: `show_id`, `title`, `type`, `date_added`, `release_year`, `rating`, `description`
- Categorical: `country`, `director`, `cast`
- Genre(s): originally in a single column `listed_in`

Missing IMDb ratings were fetched via the OMDb API (14). Table 1 provides a summary.

Characteristic	Value
Total Titles	~8,797
Movies vs. TV Shows (%)	~70% Movies, 30% TV
Mean IMDb Rating	~6.4
Earliest Release Year	1925
Latest Release Year	2021

Table 1: *Basic Netflix Dataset Characteristics (example).*

3.2 Data Cleaning & Representation

We replaced missing director, cast, or country with “Unknown” and used the mode to fill missing rating. We dropped rows missing date_added. For duration, we extracted duration_minutes if the item is a Movie (e.g. “90 min”), or duration_seasons if the item is a TV Show (e.g. “2 Seasons”). Any remaining null IMDb ratings were replaced by the median (6.4).

3.3 Data Visualization

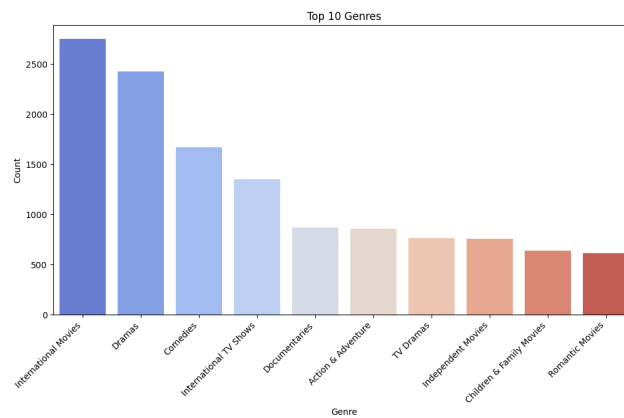


Figure 1: *Top 10 Most Popular Genres*

This bar plot clearly shows that International Movies, Dramas, and Comedies make up the bulk of Netflix’s content library, with International Movies significantly leading the count. This suggests that Netflix offers a large volume of content from different countries, and genres like drama and comedy are particularly dominant in its catalog.




Figure 2: *Distribution of IMDb Ratings*

This histogram visualizes the distribution of IMDb ratings for the movies and TV shows in our dataset. The majority of titles have ratings between 5.5 and 7.5, with a noticeable peak around 6.5 to 7.0, indicating that most content on Netflix is rated in the mid-to-high range. The distribution follows a slightly right-skewed pattern, with fewer titles receiving very low or very high ratings. The overlaid density plot further highlights the central tendency of the dataset, helping us understand how IMDb ratings are distributed across different titles.




Figure 3: *Proportion of Movies and Tv Shows*

Another important feature of the dataset is the content type—whether the content is a movie or a TV show. By examining the proportion of movies and TV shows in the dataset, we can better understand the balance of content that Netflix offers to its users. As illustrated by the pie chart, 69.6 percent of the dataset consists of movies, while 30.4 percent comprises TV shows.




Figure 4: *Proportion of Movies and Tv Shows*

This histogram shows that the majority of Netflix’s content has been released since 2010, with a dramatic rise in content production from 2020 onward. This trend is indicative of the significant increase in Netflix’s library during the streaming boom of recent years, as well as the company’s strategic push to produce more content in-house.

3.4 Feature Engineering

We split each item’s genre string, then applied a `MultiLabelBinarizer` to produce a binary vector for each possible genre. We also scaled `release_year` to $[0,1]$ using `MinMaxScaler` and stored it as `release_year_scaled`. We similarly scaled IMDb ratings to get `normalized_imdb`. The final “feature matrix” for each item is 47-dimensional (or so), containing `release_year_scaled`, `duration_minutes`, `normalized_imdb`, and binary genre flags.

4 Methods: Multi-Method Pipeline

4.1 Content-Based Similarity (Cosine)

We represent each user by the weighted average of vectors for the items they “like,” weighting by each item’s IMDb rating. Then for each candidate item, we compute:

$$\text{cosine_similarity} = \frac{\langle \text{user_profile}, \text{item_vector} \rangle}{\|\text{user_profile}\| \|\text{item_vector}\|}.$$

To incorporate quality, we define an *adjusted_similarity*:

$$\text{adjusted_similarity} = \alpha \times \text{cosine_similarity} + (1 - \alpha) \times \text{normalized_imdb}.$$

We typically set $\alpha = 0.7$. Items the user explicitly disliked or that fail the user’s IMDb threshold are excluded, and the top 10 are returned.

4.2 kNN-Based Retrieval

Rather than compute similarity for *every* item, we also do a local search using `NearestNeighbors` with cosine distance.

1. Fit kNN on the entire feature matrix (8,797 items).
2. Query the user profile vector to find, say, the top-20 neighbors.
3. Recompute the same adjusted similarity formula for these neighbors, filter out disliked/low-IMDb items, then pick the top 10.

This can yield a slightly different set, sometimes more variety.

4.3 Naïve Bayes Classification

We define a “synthetic label” `like = 1` if `imdb_rating` ≥ 7.5 else 0. We exclude user-labeled items from training (to avoid data leakage), then do an 80/20 split on the remaining items. Using *GaussianNB*, we get a probability `nb_score` $\in [0, 1]$ for each item. We reapply the trained model on the *full* dataset to produce an `nb_score` for all items.

4.4 K-Means Clustering

We tried $k = 3, 4, \dots, 9$ with silhouette scoring. Suppose $k = 4$ was selected. Each item is assigned a cluster. We compute:

$$\text{cluster_score} = \frac{\# \text{ of liked items in that cluster}}{\# \text{ of liked items total}}.$$

Hence, items in the same cluster as the user’s liked items might get up to 1.0 if that cluster is entirely user-liked. This `cluster_score` helps promote items from the user’s “favorite clusters.”

4.5 Association Rule Mining (Apriori)

We treat each item’s genres as a transaction. Using `min_support=0.05`, `min_confidence=0.6`, we find rules like (`{Independent Movies} → {Dramas}`) with confidence ~ 0.78 , lift ~ 2.82 . If an item’s genre set triggers a discovered rule, we may assign a small `association_bonus=0.05`.

4.6 Ensemble Score

To unify all signals, we define:

$$\text{ensemble_score} = w_{adj} \text{adjusted_similarity} + w_{nb} \text{nb_score} + w_{cl} \text{cluster_score} + \text{association_bonus}.$$

We exclude disliked items, sort by `ensemble_score`, and show the top 10. The weights (w_{adj} , w_{nb} , w_{cl}) can be chosen (e.g., sum to 1.0) to balance quality vs. diversity.

5 Results and Discussion

5.1 Google Form Preferences Results

Category	Selections
Liked Movies	Don, Dragon Quest Your Story, Prom Night, Miss Virginia
Disliked Movies	Twogether, Brother
Preferred Genres	International Movies, Romantic Movies, Comedies
IMDb Minimum Rating	7.2
Content Preference	Movies Only

Table 2: *User Preferences for Movie Recommendations.*

Here are the results that were extracted from the Google Form that the user filled out. These results were with the help of Google API’s: Form and Sheets. The app’s script was used to extract random movies and tv shows from the data set so when a user fills out a form they will see a set of different titles each time.

5.2 Content-Based Retrival (Cosine Similarity)

Rank	Title	Type	IMDb Rating	Adjusted Similarity
1	Dick Johnson Is Dead	Movie	7.4	0.920910
2	Jaws	Movie	8.1	0.947406
3	Training Day	Movie	7.8	0.936165
4	InuYasha the Movie 2: The Castle Beyond the Looking Glass	Movie	7.3	0.917307
5	InuYasha the Movie 3: Swords of an Honorable Ruler	Movie	7.6	0.928512
6	InuYasha the Movie 4: Fire on the Mystic Island	Movie	7.2	0.913403
7	InuYasha the Movie: Affections Touching Across Time	Movie	7.2	0.913581
8	Schumacher	Movie	7.4	0.921120
9	Omo Ghetto: the Saga	Movie	7.6	0.928715
10	If I Leave Here Tomorrow: A Film About Lynnyrd Skynyrd	Movie	7.8	0.935915

Table 3: *Top 10 Recommendations Based on Cosine Similarity.*

5.3 kNN-Based Retrieval

Rank	Title	Type	IMDb Rating	Adjusted Similarity
1	Django Unchained	Movie	8.5	0.962482
2	Inglourious Basterds	Movie	8.4	0.958737
3	Casino Royale	Movie	8.0	0.943736
4	GoldenEye	Movie	7.2	0.913734

Table 4: *Top kNN-Based Recommendations.*

5.4 Naïve Bayes Performance

Excluding user-labeled items from training, we got $\sim 84\%$ accuracy, ~ 0.54 precision, ~ 0.59 recall, and 0.57 F1. Despite moderate performance on the synthetic “IMDb7.5” label, *nb_score* still offers a useful dimension for boosting items likely to be user “likes.”

Metric	Score
Accuracy	0.84
Precision	0.54
Recall	0.59
F1 Score	0.57

Table 5: *Naïve Bayes Classifier Performance Metrics.*

5.5 Clustering Outcomes

5.5.1 Elbow Method for Optimal k




Figure 5: *Elbow Method for Optimal K Selection in K-Means Clustering*

5.5.2 Cluster Distribution and Silhouette Score

Number of Clusters (k)	Cluster Distribution	Silhouette Score
3	[2915, 4155, 1727]	0.665
4	[2807, 3650, 1326, 1014]	0.659
5	[2796, 1787, 571, 810, 2833]	0.649
6	[2779, 2045, 339, 667, 1778, 1189]	0.625
7	[2766, 2253, 656, 550, 1082, 1327, 163]	0.613
8	[2689, 1671, 710, 632, 201, 1137, 254, 1503]	0.610
9	[2689, 1573, 722, 632, 281, 1103, 254, 1503, 40]	0.610

Table 6: *K-Means Clustering Results for Different k Values.*

We found $k = 4$ gave a silhouette of around 0.659. Table 7 shows the cluster distribution. Each cluster had a distinct set of common genres and a different average IMDb rating.

Cluster	Size	Avg IMDb	Top Genres
0	2,807	6.92	{International TV Shows, TV Dramas, TV Comedies}
1	3,650	6.16	{International Movies, Dramas, Comedies}
2	1,326	6.66	{International Movies, Dramas}
3	1,014	6.42	{Documentaries, Stand-up Comedy}

Table 7: *Cluster distribution and top genres at $k = 4$.*

5.6 Final Recommendations

After computing ensemble scores, the top items typically balanced:

- `adjusted_similarity` (cosine with user profile, plus IMDb weighting),
- `nb_score` (Naïve Bayes probability),
- `cluster_score` (user’s liked cluster fraction),
- `association_bonus` (0.05 if the item matches strong genre rules).

Table 8 illustrates sample final recommendations for a user with certain likes/dislikes.

Rank	Title	Type	IMDb Rating	Ensemble Score
1	Monty Python and the Holy Grail	Movie	8.2	0.904121
2	Monty Python’s Life of Brian	Movie	8.0	0.901261
3	Cat on a Hot Tin Roof	Movie	7.9	0.899925
4	Ferris Bueller’s Day Off	Movie	7.8	0.898450
5	Willy Wonka and the Chocolate Factory	Movie	7.8	0.898439
6	Bonnie and Clyde	Movie	7.7	0.897060
7	Rebel Without a Cause	Movie	7.6	0.895620
8	Enter the Dragon	Movie	7.6	0.895600
9	Dark Waters	Movie	7.6	0.895591
10	Forbidden Planet	Movie	7.5	0.894122

Table 8: *Final Top 10 Movie Recommendations Based on Ensemble Scoring.*

The user-labeled disliked items are excluded, and we see comedic, older classics that still align with user interest.

5.7 Limitations

We rely solely on item content (no user-user collaborative data). The Naïve Bayes “like=IMDb7.5” is synthetic, potentially missing real user preferences. Some lift-laden association rules might not reflect actual user taste. But this

system does illustrate how combining multiple content-based signals yields more robust coverage than a single approach.

6 Conclusion and Future Work

We presented a multi-method content-based recommender system for Netflix data, integrating:

- Cosine-based user profile matching,
- kNN local retrieval,
- Naïve Bayes synthetic classification,
- K-Means cluster preference,
- Apriori association rules.

Our pipeline overcame many typical single-method limitations, fostering more diverse, user-aligned recommendations. While successful in bridging item-based signals, we have not yet incorporated collaborative filtering or real user rating data. We also used synthetic labels for the Naïve Bayes step.

Future directions include collecting actual user ratings, exploring dynamic weighting for association rules, or a hybrid with partial user-user data. Overall, this project demonstrates that a synergy of multiple content-based techniques can handle large catalogs like Netflix and yield personalized results with minimal cold start overhead.

References

- [1] D. Goldberg, D. Nichols, B. M. Oki, & D. Terry. Using collaborative filtering to weave an information tapestry. *Communications of the ACM*, 35(12):61–70, 1992.
- [2] P. Resnick, N. Iacovou, M. Suchak, P. Bergstrom, & J. Riedl. GroupLens: An open architecture for collaborative filtering of netnews. In *Proceedings of the 1994 ACM conference on Computer supported cooperative work*, pp. 175–186, 1994.
- [3] A. I. Schein, A. Popescul, L. H. Ungar, & D. M. Pennock. Methods and metrics for cold-start recommendations. In *Proceedings of the 25th annual international ACM SIGIR conference on Research and development in information retrieval*, pp. 253–260, 2002.

- [4] R. J. Mooney & L. Roy. Content-based book recommending using learning for text categorization. In *Proceedings of the fifth ACM conference on Digital libraries*, pp. 195–204, 2000.
- [5] M. Pazzani & D. Billsus. Content-based recommendation systems. In *The Adaptive Web*, pages 325–341. Springer, 2007.
- [6] R. Burke. Hybrid recommender systems: Survey and experiments. *User Modeling and User-Adapted Interaction*, 12(4):331–370, 2002.
- [7] G. Adomavicius & A. Tuzhilin. Toward the next generation of recommender systems: a survey of the state-of-the-art and possible extensions. *IEEE Transactions on Knowledge and Data Engineering*, 17(6):734–749, 2005.
- [8] C. A. Gómez-Urbe & N. Hunt. The Netflix recommender system: Algorithms, business value, and innovation. *ACM Transactions on Management Information Systems*, 6(4):1–19, 2015.
- [9] E. Alpaydm. *Introduction to Machine Learning* (4th ed.). MIT Press, 2020.
- [10] A. Das, M. Datar, A. Garg, & S. Rajaram. Google news personalization: scalable online collaborative filtering. In *Proceedings of the 16th International Conference on World Wide Web*, pp. 271–280, 2007.
- [11] R. Agrawal & R. Srikant. Fast algorithms for mining association rules. In *Proc. 20th int. conf. very large data bases (VLDB)*, pages 487–499, 1994.
- [12] B. Mobasher, H. Dai, T. Luo, & M. Nakagawa. Effective personalization based on association rule discovery from web usage data. In *Proceedings of the 3rd international workshop on Web information and data management*, pp. 9–15, 2001.
- [13] Chirag9073. Netflix Data Analysis. In *Kaggle.com*, Kaggle, June 3, 2020. Available at: <https://www.kaggle.com/code/chirag9073/netflix-data-analysis/notebook>.
- [14] OMDb API. The Open Movie Database. In *Omdbapi.com*, 2000. Available at: <https://www.omdbapi.com/>.

7 Appendix

```
1
2 # Import required libraries and mount Google Drive
3
4 import os, time, math, io, requests
5 from concurrent.futures import ThreadPoolExecutor
6
7 # Data handling and numerical computation
8 import numpy as np
9 import pandas as pd
10 import scipy.stats as ss
11
12 # Machine learning and preprocessing
13 from sklearn.preprocessing import MultiLabelBinarizer,
14     ↪ MinMaxScaler
15 from sklearn.metrics.pairwise import cosine_similarity
16 from sklearn.neighbors import NearestNeighbors
17 from sklearn.naive_bayes import GaussianNB
18 from sklearn.cluster import KMeans
19 from sklearn.model_selection import train_test_split
20 from sklearn.metrics import accuracy_score,
21     ↪ precision_score, recall_score, f1_score
22 from sklearn.metrics import silhouette_score
23 from mlxtend.frequent_patterns import apriori,
24     ↪ association_rules
25
26 # Visualization libraries
27 import seaborn as sns
28 import matplotlib.pyplot as plt
29
30 # Google API libraries
31 !pip install --upgrade gspread google-auth
32     ↪ google-auth-oauthlib google-auth-httpplib2
33     ↪ google-api-python-client
34
35 import gspread
36 from google.oauth2.service_account import Credentials
37
38 # Progress bar utility
39 from tqdm import tqdm
40
41 # Mount Google Drive (for Colab)
42 from google.colab import drive
43 drive.mount('/content/drive', force_remount=True)
44
```



```

39 # working directory project folder in Google Drive
40 pathname = '/content/drive/MyDrive/DataMiningTeam/'
41 os.chdir(pathname)
42
43
44 # Load the raw Netflix dataset and perform data cleaning
45
46 filename = "netflix_database.xlsx"
47 dt = pd.read_excel(filename)
48
49 print("Initial dataset preview:")
50 display(dt.head())
51
52 # Fill missing values for director, cast, and country
53 dt[['director', 'cast', 'country']] = dt[['director',
54     ↳ 'cast', 'country']].fillna("Unknown")
55
56 # Fill missing ratings with the most common rating (mode)
57 mode_rating = dt['rating'].mode()[0]
58 dt['rating'] = dt['rating'].fillna(mode_rating)
59
60 # Rename 'listed_in' to 'genre'
61 dt.rename(columns={'listed_in': 'genre'}, inplace=True)
62
63 # Drop rows with missing 'date_added' and extract day,
64     ↳ month, year
65 dt = dt.dropna(subset=['date_added']).copy()
66 dt['day_added'] = dt['date_added'].dt.day.astype(int)
67 dt['month_added'] = dt['date_added'].dt.month.astype(int)
68 dt['year_added'] = dt['date_added'].dt.year.astype(int)
69
70 # Fill missing duration separately for TV Shows and Movies
71     ↳ using mode
72 tv_mode = dt.loc[dt['type'] == 'TV Show',
73     ↳ 'duration'].mode()[0]
74 mv_mode = dt.loc[dt['type'] == 'Movie',
75     ↳ 'duration'].mode()[0]
76 dt.loc[dt['type'] == 'TV Show', 'duration'] =
77     ↳ dt.loc[dt['type'] == 'TV Show',
78     ↳ 'duration'].fillna(tv_mode)
79 dt.loc[dt['type'] == 'Movie', 'duration'] =
80     ↳ dt.loc[dt['type'] == 'Movie',
81     ↳ 'duration'].fillna(mv_mode)
82
83 # Extract numerical duration:

```

```

75 # For movies: extract minutes; for TV shows: extract number
    ↪ of seasons.
76 dt['duration_minutes'] = dt['duration'].apply(lambda x:
    ↪ float(x.split()[0]) if 'min' in str(x) else np.nan)
77 dt['duration_seasons'] = dt['duration'].apply(lambda x:
    ↪ float(x.split()[0]) if 'Season' in str(x) else np.nan)
78 dt.drop(columns=['duration'], inplace=True)
79
80 print("Missing values after cleaning:")
81 print(dt.isnull().sum())
82 print("\nDataset preview after cleaning:")
83 display(dt.head())
84
85
86 # Fetch missing IMDb ratings using the OMDb API with
    ↪ multithreading
87
88 OMDb_API_KEY = "17bd2674"
89 OMDb_URL = "http://www.omdbapi.com/"
90
91 def get_imdb_rating(title):
92     params = {"t": title, "apikey": OMDb_API_KEY}
93     try:
94         response = requests.get(OMDb_URL, params=params,
    ↪ timeout=5)
95         data = response.json()
96         if response.status_code == 200 and
    ↪ data.get("Response") == "True":
97             return data.get("imdbRating", "N/A")
98     except Exception as e:
99         print(f"Error fetching rating for {title}: {e}")
100     return None
101
102 # Ensure 'imdb_rating' exists
103 if "imdb_rating" not in dt.columns:
104     dt["imdb_rating"] = None
105
106 titles_to_fetch = dt[(dt["imdb_rating"].isna()) |
    ↪ (dt["imdb_rating"] == "N/A")]["title"].tolist()
107
108 with ThreadPoolExecutor(max_workers=20) as executor:
109     results = list(tqdm(executor.map(get_imdb_rating,
    ↪ titles_to_fetch), total=len(titles_to_fetch)))
110
111 for title, imdb in zip(titles_to_fetch, results):

```

```

112         dt.loc[dt["title"] == title, "imdb_rating"] = imdb
113
114     # Convert to numeric and fill any missing values with the
115     ↪ median
116     dt["imdb_rating"] = pd.to_numeric(dt["imdb_rating"],
117     ↪ errors='coerce')
118     median_rating = dt["imdb_rating"].median()
119     dt["imdb_rating"].fillna(median_rating, inplace=True)
120
121     file_path = "/content/drive/My
122     ↪ Drive/DataMiningTeam/netflix_database_cleaned.xlsx"
123     dt.to_excel(file_path, index=False)
124     print(f"IMDb Ratings Fetching Complete! Data saved to:
125     ↪ {file_path}")
126
127
128 # Cell 4: Visualize dataset characteristics
129
130 def plot_visualizations(dt):
131     plt.figure(figsize=(6, 4))
132     sns.countplot(x='type', data=dt)
133     plt.title('Count of Movies vs TV Shows')
134     plt.show()
135
136     plt.figure(figsize=(10, 5))
137     sns.histplot(dt["imdb_rating"].dropna(), bins=20,
138     ↪ kde=True)
139     plt.title("Distribution of IMDb Ratings")
140     plt.xlabel("IMDb Rating")
141     plt.ylabel("Count")
142     plt.show()
143
144     plt.figure(figsize=(10, 5))
145     sns.countplot(x='rating', data=dt,
146     ↪ order=dt['rating'].value_counts().index)
147     plt.title('Content Rating Distribution')
148     plt.xticks(rotation=45)
149     plt.show()
150
151     plt.figure(figsize=(8, 5))
152     sns.histplot(x='release_year', data=dt, bins=30)
153     plt.title('Release Year Distribution')
154     plt.show()
155
156     plt.figure(figsize=(8, 5))

```

```

151     sns.histplot(x='duration_minutes', data=dt[dt['type']
    ↪     == 'Movie'])
152     plt.title('Movie Duration (Minutes)')
153     plt.show()
154
155     plt.figure(figsize=(8, 5))
156     sns.histplot(x='duration_seasons', data=dt[dt['type']
    ↪     == 'TV Show'], bins=range(1,15))
157     plt.title('TV Show Seasons Distribution')
158     plt.show()
159
160     dt["genre_list"] = dt["genre"].fillna("").apply(lambda
    ↪     x: x.split(", "))
161     all_genres = [genre for sublist in dt["genre_list"] for
    ↪     genre in sublist if genre]
162     top_genres =
    ↪     pd.Series(all_genres).value_counts().head(10)
163     plt.figure(figsize=(12, 6))
164     sns.barplot(x=top_genres.index, y=top_genres.values,
    ↪     palette="coolwarm")
165     plt.xticks(rotation=45, ha="right")
166     plt.title("Top 10 Genres")
167     plt.xlabel("Genre")
168     plt.ylabel("Count")
169     plt.show()
170
171 plot_visualizations(dt)
172
173
174 # engineered features
175
176 # Created a genre list from the 'genre' column
177 dt['genre_list'] = dt['genre'].fillna("").apply(lambda x:
    ↪     x.split(", "))
178
179 # One-hot encode genres
180 mlb = MultiLabelBinarizer()
181 genre_encoded = mlb.fit_transform(dt['genre_list'])
182 genre_df = pd.DataFrame(genre_encoded,
    ↪     columns=mlb.classes_, index=dt.index)
183
184 # Scale release_year
185 scaler = MinMaxScaler()
186 dt['release_year_scaled'] =
    ↪     scaler.fit_transform(dt[['release_year']].fillna(dt['release_year'].min()))

```

```

187
188 # final feature-engineered DataFrame
189 dt_final = pd.concat([
190     dt[['show_id', 'title', 'type', 'release_year_scaled',
191         ↳ 'duration_minutes', 'duration_seasons']],
192     genre_df
193 ], axis=1).copy()
194
195 # Add IMDb rating
196 dt_final["imdb_rating"] = dt["imdb_rating"]
197
198 # Save engineered dataset
199 file_path = "/content/drive/My
200 ↳ Drive/DataMiningTeam/netflix_feature_engineered.xlsx"
201 dt_final.to_excel(file_path, index=False)
202 print(f"Feature Engineered Dataset Saved: {file_path}")
203
204 print("Shape of dt_final:", dt_final.shape)
205 display(dt_final.head())
206
207 # Construct the numerical feature matrix for similarity
208 ↳ calculations
209 # and normalize IMDb ratings in one cell.
210
211 # First, update dt_final with normalized IMDb ratings.
212 dt_final["imdb_rating"] =
213 ↳ pd.to_numeric(dt_final["imdb_rating"], errors='coerce')
214 median_imdb = dt_final["imdb_rating"].median()
215 dt_final["imdb_rating"] =
216 ↳ dt_final["imdb_rating"].fillna(median_imdb)
217
218 scaler = MinMaxScaler()
219 dt_final["normalized_imdb"] =
220 ↳ scaler.fit_transform(dt_final[["imdb_rating"]])
221
222 # Now, drop non-numeric columns and fill any remaining
223 ↳ missing values.
224 feature_columns = dt_final.drop(columns=['show_id',
225     ↳ 'title', 'type'], errors='ignore').fillna(0)
226 feature_matrix = feature_columns.to_numpy()
227
228 print("Feature Matrix Created!")
229 print("Shape:", feature_matrix.shape)
230 print("Sample rows:\n", feature_matrix[:1])

```

```

224
225
226 # Load user preferences via Google Sheets
227
228 # Load user preferences via Google Sheets and parse them
    ↳ into usable Python objects.
229
230 SERVICE_ACCOUNT_FILE = "/content/drive/My
    ↳ Drive/DataMiningTeam/service_account.json"
231
232 creds = Credentials.from_service_account_file(
233     SERVICE_ACCOUNT_FILE,
234     scopes=["https://spreadsheets.google.com/feeds",
    ↳ "https://www.googleapis.com/auth/drive"]
235 )
236 gc = gspread.authorize(creds)
237 print("Google Sheets Authentication Successful!")
238
239 SPREADSHEET_URL =
    ↳ "https://docs.google.com/spreadsheets/d/1lCOs6jCFzElSMi78x39_VZtGqhVOpilp
240 spreadsheet = gc.open_by_url(SPREADSHEET_URL)
241 worksheet = spreadsheet.sheet1
242 data = worksheet.get_all_values()
243 latest_response = pd.DataFrame([data[-1]], columns=data[0])
244 print("User History Loaded!")
245
246 # Parse user responses
247 latest_response.columns =
    ↳ latest_response.columns.str.strip()
248
249 liked_movies = latest_response["Liked Movies/TV
    ↳ Shows"].values[0] if "Liked Movies/TV Shows" in
    ↳ latest_response else ""
250 disliked_movies = latest_response["Disliked Movies/TV
    ↳ Shows"].values[0] if "Disliked Movies/TV Shows" in
    ↳ latest_response else ""
251 preferred_genres = latest_response["Preferred
    ↳ Genres"].values[0] if "Preferred Genres" in
    ↳ latest_response else ""
252 imdb_min_rating = latest_response["IMDb Rating
    ↳ Preference"].values[0] if "IMDb Rating Preference" in
    ↳ latest_response else "No preference"
253 content_preference = latest_response["Do you prefer Movies
    ↳ or TV Shows?"].values[0] if "Do you prefer Movies or TV
    ↳ Shows?" in latest_response else "No preference"

```

```

254 liked_movies = liked_movies.split(", ") if liked_movies
255     ↪ else []
256 disliked_movies = disliked_movies.split(", ") if
257     ↪ disliked_movies else []
258 preferred_genres = preferred_genres.split(", ") if
259     ↪ preferred_genres else []
260
261 if imdb_min_rating not in ["No preference", ""]:
262     imdb_min_rating = float(imdb_min_rating.rstrip("+"))
263 else:
264     imdb_min_rating = 0
265
266 print(f"Liked Movies: {liked_movies}")
267 print(f"Disliked Movies: {disliked_movies}")
268 print(f"Preferred Genres: {preferred_genres}")
269 print(f"IMDb Minimum Rating: {imdb_min_rating}")
270 print(f"Content Preference: {content_preference}")
271
272
273 # Compute a weighted user profile based on liked movies and
274 ↪ calculate cosine similarity.
275
276 # Find indices of movies the user likes.
277 liked_indices =
278     ↪ dt_final[dt_final['title'].isin(liked_movies)].index.to_numpy()
279
280 # If no liked movies but preferred genres exist, use a
281 ↪ fallback: select top 3 movies in those genres.
282 if len(liked_indices) == 0 and preferred_genres:
283     genre_filtered =
284         ↪ dt_final[dt_final[preferred_genres].sum(axis=1) >
285         ↪ 0]
286     liked_movies = genre_filtered['title'].head(3).tolist()
287     liked_indices =
288         ↪ dt_final[dt_final['title'].isin(liked_movies)].index.to_numpy()
289
290 # Compute user profile: weighted average of feature vectors
291 ↪ for liked movies.
292 if len(liked_indices) > 0:
293     weights = dt_final.iloc[liked_indices]['imdb_rating'] /
294         ↪ dt_final.iloc[liked_indices]['imdb_rating'].sum()

```

```

287     user_profile =
        ↳ np.average(feature_matrix[liked_indices], axis=0,
        ↳ weights=weights).reshape(1, -1)
288 else:
289     # Fallback to using the average feature vector across
        ↳ all items if no liked movies.
290     user_profile = np.mean(feature_matrix,
        ↳ axis=0).reshape(1, -1)
291
292 # Compute cosine similarity between the user profile and
        ↳ all items.
293 similarities = cosine_similarity(user_profile,
        ↳ feature_matrix)[0]
294 dt_final['similarity_score'] = similarities
295
296 # Compute adjusted similarity using a weighted sum (alpha
        ↳ for similarity, 1-alpha for normalized IMDb rating).
297 alpha = 0.7
298 dt_final["adjusted_similarity"] = alpha *
        ↳ dt_final["similarity_score"] + (1 - alpha) *
        ↳ dt_final["normalized_imdb"]
299
300 print("User Profile & Cosine Similarity Computed!")
301
302
303
304
305 # Filter recommendations based on user preferences and
        ↳ display final content-based recommendations.
306
307 # Exclude items the user already liked or disliked.
308 recommended =
        ↳ dt_final[~dt_final['title'].isin(liked_movies)]
309 recommended =
        ↳ recommended[~recommended['title'].isin(disliked_movies)]
310 # Apply the IMDb rating filter.
311 recommended = recommended[recommended["imdb_rating"] >=
        ↳ imdb_min_rating]
312
313 # Fallback: If fewer than 10 recommendations are found,
        ↳ expand search using preferred genres.
314 if len(recommended) < 10:
315     print("Less than 10 strong matches found. Expanding
        ↳ search with fallback recommendations...")
316     if preferred_genres:

```



```

317         fallback =
            ↳ dt_final[dt_final[preferred_genres].sum(axis=1)
            ↳ > 0]
318         fallback = fallback.sort_values(by=["imdb_rating",
            ↳ "similarity_score"], ascending=[False, False])
319         recommended = pd.concat([recommended,
            ↳ fallback]).drop_duplicates().head(10)
320     else:
321         recommended = recommended.head(10)
322
323     # If the user prefers a specific type, split
    ↳ recommendations accordingly.
324     recommended_movies = recommended[recommended["type"] ==
    ↳ "Movie"]
325     recommended_shows = recommended[recommended["type"] == "TV
    ↳ Show"]
326
327     if content_preference == "Both":
328         top_movies = recommended_movies.head(5)
329         top_shows = recommended_shows.head(5)
330         final_recommendations = pd.concat([top_movies,
            ↳ top_shows]).sort_values(by="adjusted_similarity",
            ↳ ascending=False)
331
332     elif content_preference == "Movies Only":
333         final_recommendations = recommended_movies.head(10)
334
335     elif content_preference == "TV Shows Only":
336         final_recommendations = recommended_shows.head(10)
337
338     else: # "No preference" or any other input
339         final_recommendations = recommended.head(10)
340
341     print("Top Recommendations:Based on Cosine Similarity")
342     display(final_recommendations[['title', 'type',
    ↳ 'imdb_rating', 'adjusted_similarity']])
343
344
345
346     # alternative recommendations using kNN.
347
348     # Initialize and fit the kNN model on the full feature
    ↳ matrix
349     knn = NearestNeighbors(n_neighbors=20, metric='cosine')
350     knn.fit(feature_matrix)

```

```

351
352 # Find the indices of the 20 nearest neighbors to the user
    ↳ profile
353 _, indices = knn.kneighbors(user_profile)
354 knn_recommendations = dt_final.iloc[indices[0]].copy()
355
356 # Compute cosine similarity for these kNN recommendations
    ↳ using their numeric features
357 # Note: Reconstruct the feature matrix for just these
    ↳ neighbors
358 neighbors_features =
    ↳ feature_columns.iloc[indices[0]].to_numpy()
359 knn_similarities = cosine_similarity(user_profile,
    ↳ neighbors_features)[0]
360 knn_recommendations['similarity_score'] = knn_similarities
361
362 # Compute adjusted similarity as a weighted sum of cosine
    ↳ similarity and normalized IMDb rating
363 alpha = 0.7 # weight for similarity; 30% weight for
    ↳ normalized IMDb rating
364 knn_recommendations["adjusted_similarity"] = alpha *
    ↳ knn_recommendations["similarity_score"] + (1 - alpha) *
    ↳ knn_recommendations["normalized_imdb"]
365
366 # Exclude items the user disliked
367 knn_recommendations =
    ↳ knn_recommendations[~knn_recommendations['title'].isin(disliked_movies)]
368
369 # Apply the minimum IMDb rating filter from user preference
370 knn_recommendations =
    ↳ knn_recommendations[knn_recommendations["imdb_rating"]
    ↳ >= imdb_min_rating]
371
372 # Filter recommendations based on content preference:
373 if content_preference == "Movies Only":
374     knn_recommendations =
        ↳ knn_recommendations[knn_recommendations["type"] ==
        ↳ "Movie"]
375 elif content_preference == "TV Shows Only":
376     knn_recommendations =
        ↳ knn_recommendations[knn_recommendations["type"] ==
        ↳ "TV Show"]
377
378

```

```

379 # Sort by adjusted similarity in descending order for final
    ↳ ranking
380 knn_recommendations =
    ↳ knn_recommendations.sort_values(by="adjusted_similarity",
    ↳ ascending=False)
381
382 print("Top kNN-Based Recommendations:")
383 display(knn_recommendations[['title', 'type',
    ↳ 'imdb_rating', 'adjusted_similarity']].head(10))
384
385
386 # Association Rule Mining on Genre Data using Apriori
387
388 # Convert genre_df to boolean type.
389 genre_df = genre_df.astype(bool)
390
391 frequent_itemsets = apriori(genre_df, min_support=0.05,
    ↳ use_colnames=True)
392 rules = association_rules(frequent_itemsets,
    ↳ metric="confidence", min_threshold=0.6)
393
394 print("Association Rules based on genres (sample):")
395 display(rules[['antecedents', 'consequents', 'support',
    ↳ 'confidence', 'lift']].head())
396
397
398 # Naïve Bayes classifier
399
400 # 1. Define 'like' using your synthetic rule
401 dt_final['like'] = (dt_final['imdb_rating'] >=
    ↳ 7.5).astype(int)
402
403 # 2. Combine user-labeled items
404 user_labeled_titles = set(liked_movies + disliked_movies)
405
406 # 3. Exclude those items from dt_final before building the
    ↳ train set
407 excluded_indices =
    ↳ dt_final[dt_final['title'].isin(user_labeled_titles)].index
408 train_dt = dt_final.drop(index=excluded_indices)
409
410 # 4. List columns we DO NOT want for model training
411 exclude_cols = [
412     'show_id', 'title', 'type', # not features

```

```

413     'like', 'cluster', 'similarity_score', # or any added
        ↳ columns
414     'adjusted_similarity', 'ensemble_score',
415     'association_bonus', 'nb_score', 'cluster_score'
416 ]
417
418 # 5. Prepare the training feature matrix
419 train_feature_cols = train_dt.drop(columns=exclude_cols,
        ↳ errors='ignore').fillna(0)
420 train_X = train_feature_cols.to_numpy()
421 train_y = train_dt['like'].values
422
423 # 6. Train/Test split
424 X_train, X_test, y_train, y_test = train_test_split(
425     train_X,
426     train_y,
427     test_size=0.2,
428     random_state=42
429 )
430
431 clf = GaussianNB()
432 clf.fit(X_train, y_train)
433
434 predictions = clf.predict(X_test)
435
436 # Evaluate the classifier
437 acc = accuracy_score(y_test, predictions)
438 prec = precision_score(y_test, predictions)
439 rec = recall_score(y_test, predictions)
440 f1 = f1_score(y_test, predictions)
441
442 print("Naïve Bayes classifier trained (excluding
        ↳ user-labeled items).")
443 print(f"Accuracy: {acc:.2f}, Precision: {prec:.2f}, Recall:
        ↳ {rec:.2f}, F1 Score: {f1:.2f}")
444
445
446 # Re-applying the model to the FULL dt_final so we can get
        ↳ nb_score for all items
447 full_feature_cols = dt_final.drop(columns=exclude_cols,
        ↳ errors='ignore').fillna(0)
448 full_X = full_feature_cols.to_numpy()
449
450 nb_proba_all = clf.predict_proba(full_X)[: , 1]
451 dt_final['nb_score'] = nb_proba_all

```

```

452
453 print("Example nb_score:")
454 display(dt_final[['title', 'nb_score']].head(10))
455
456
457 # Elbow
458
459 sse = []
460 for k in range(2, 11):
461     kmeans = KMeans(n_clusters=k, random_state=42)
462     kmeans.fit(feature_matrix)
463     sse.append(kmeans.inertia_)
464
465 plt.figure(figsize=(8, 5))
466 plt.plot(range(2, 11), sse, marker='o')
467 plt.title("Elbow Method for Optimal k")
468 plt.xlabel("Number of Clusters (k)")
469 plt.ylabel("Sum of Squared Errors (SSE)")
470 plt.show()
471
472
473 # Apply K-Means clustering to group items
474
475
476
477 print("Trying different k values for K-Means:")
478 for k in range(3, 10):
479     kmeans_temp = KMeans(n_clusters=k, random_state=42)
480     clusters_temp = kmeans_temp.fit_predict(feature_matrix)
481     sil_score = silhouette_score(feature_matrix,
482     ↪ clusters_temp)
483     print(f"k={k}: Cluster Distribution:
484     ↪ {np.bincount(clusters_temp)} | Silhouette Score:
485     ↪ {sil_score:.3f}")
486
487
488 # Set optimal k
489 k_optimal = 4
490 kmeans = KMeans(n_clusters=k_optimal, random_state=42)
491 clusters = kmeans.fit_predict(feature_matrix)
492 dt_final['cluster'] = clusters
493
494 print("K-Means clustering complete. Cluster distribution:")
495 print(dt_final['cluster'].value_counts())
496
497
498
499

```

```

494 # Silhouette Bar Chart
495
496 # We'll reuse the range(3, 10) and store silhouette scores
    ↳ in a list for plotting
497 k_values = list(range(3, 10))
498 silhouette_scores = []
499
500 print("\nSilhouette Bar Chart for K-Means:")
501 for k in k_values:
502     kmeans_temp = KMeans(n_clusters=k, random_state=42)
503     clusters_temp = kmeans_temp.fit_predict(feature_matrix)
504     sil_score = silhouette_score(feature_matrix,
    ↳ clusters_temp)
505     silhouette_scores.append(sil_score)
506
507 plt.figure(figsize=(7, 5))
508 plt.bar(k_values, silhouette_scores, color='skyblue')
509 plt.title("Silhouette Scores for Different k")
510 plt.xlabel("Number of Clusters (k)")
511 plt.ylabel("Silhouette Score")
512 for i, score in enumerate(silhouette_scores):
513     plt.text(k_values[i], score, f"{score:.3f}",
    ↳ ha='center', va='bottom', fontsize=9)
514 plt.ylim(0, max(silhouette_scores) + 0.05)
515 plt.show()
516
517
518 # Cluster Profiling
519
520 # Suppose you want to see the average IMDb rating,
    ↳ release_year_scaled,
521 # and the top 5 genres in each cluster.
522
523 genre_cols = genre_df.columns # from your one-hot-encoded
    ↳ df
524 unique_clusters = sorted(dt_final['cluster'].unique())
525
526 print("\n=== Cluster Profiling ===")
527 for c in unique_clusters:
528     cluster_data = dt_final[dt_final['cluster'] == c]
529     print(f"\n--- Cluster {c} ---")
530     print(f"Number of items: {len(cluster_data)}")
531     print(f"Average IMDb rating:
    ↳ {cluster_data['imdb_rating'].mean():.2f}")

```

```

532     print(f"Average release_year_scaled:
    ↪     {cluster_data['release_year_scaled'].mean():.2f}")
533
534     # Summation over genre columns to see which genres are
    ↪     most common
535     cluster_genre_sum =
    ↪     cluster_data[genre_cols].sum().sort_values(ascending=False).head(5)
536     print("Top 5 genres:")
537     display(cluster_genre_sum)
538
539
540     # Compute an ensemble recommendation score combining
    ↪     multiple methods
541
542     # Ensure 'normalized_imdb' exists; recompute if missing
543     if "normalized_imdb" not in dt_final.columns:
544         print(" 'normalized_imdb' column missing.
    ↪         Recomputing...")
545         dt_final["imdb_rating"] =
    ↪         pd.to_numeric(dt_final["imdb_rating"],
    ↪         errors='coerce')
546         median_imdb = dt_final["imdb_rating"].median()
547         dt_final["imdb_rating"] =
    ↪         dt_final["imdb_rating"].fillna(median_imdb)
548         scaler = MinMaxScaler()
549         dt_final["normalized_imdb"] =
    ↪         scaler.fit_transform(dt_final[["imdb_rating"]])
550
551     # Ensure 'adjusted_similarity' exists; recompute if missing
552     if "adjusted_similarity" not in dt_final.columns:
553         print(" 'adjusted_similarity' missing. Recomputing...")
554         similarities = cosine_similarity(user_profile,
    ↪         feature_matrix)[0]
555         dt_final['similarity_score'] = similarities
556         alpha = 0.7 # weight for similarity
557         dt_final["adjusted_similarity"] = alpha *
    ↪         dt_final["similarity_score"] + (1 - alpha) *
    ↪         dt_final["normalized_imdb"]
558         dt_final["adjusted_similarity"] =
    ↪         dt_final["adjusted_similarity"].fillna(0)
559
560
561     # Get Naïve Bayes classifier probability for the positive
    ↪     (like) class
562     nb_proba = clf.predict_proba(feature_matrix)[: , 1]

```

```

563 dt_final['nb_score'] = nb_proba
564
565 # Compute fraction of liked movies in each cluster as a
    ↳ measure of cluster preference
566 liked_clusters =
    ↳ dt_final.loc[dt_final['title'].isin(liked_movies),
    ↳ 'cluster']
567 if len(liked_clusters) > 0:
568     cluster_pref =
        ↳ liked_clusters.value_counts(normalize=True).to_dict()
569 else:
570     cluster_pref = {}
571 dt_final['cluster_score'] =
    ↳ dt_final['cluster'].apply(lambda x: cluster_pref.get(x,
    ↳ 0))
572
573 # Adjust ensemble score weights: Here, we use 0.4 for
    ↳ adjusted similarity, 0.4 for NB score, and 0.2 for
    ↳ cluster score.
574 # (tweak these weights to get a better balance of quality
    ↳ vs. diversity.)
575 total = 0.4 + 0.4 + 0.5
576
577 w_adj = 0.5/total
578 w_nb_score = 0.4/total
579 w_cluster_score = 0.4/total
580
581
582
583 dt_final['ensemble_score'] = (w_adj *
    ↳ dt_final['adjusted_similarity'] +
584     w_nb_score *
        ↳ dt_final['nb_score'] +
585     w_cluster_score *
        ↳ dt_final['cluster_score'])
586
587 # association rule
588 associated_titles = []
589
590 dt_final['association_bonus'] =
    ↳ dt_final['title'].apply(lambda t: 0.05 if t in
    ↳ associated_titles else 0)
591 dt_final['ensemble_score'] += dt_final['association_bonus']
592

```



```
593 | # Filter out liked and disliked items from ensemble
    | ↳ recommendations
594 | ensemble_recs =
    | ↳ dt_final[~dt_final['title'].isin(liked_movies)]
595 | ensemble_recs =
    | ↳ ensemble_recs[~ensemble_recs['title'].isin(disliked_movies)]
596 | ensemble_recs =
    | ↳ ensemble_recs.sort_values(by='ensemble_score',
    | ↳ ascending=False)
597 |
598 | print(" Final Recommendations:")
599 | display(ensemble_recs[['title', 'type', 'imdb_rating',
    | ↳ 'ensemble_score']].head(10))
600 |
601 |
602 |
```